

Cache Advanced Features

Before the advanced features for the cache can be implemented, our cache design from MP3 must be updated to respond in 1 cycle. In the MP3 implementation, the array register output values are only updated by the end of the clock cycle. This means, for example, that the 'hit' signal takes the whole cycle to generate, making a 1-hit design impossible. In order to adjust this, the arrays' dataout value is generated in a combinational block. The extra state in the state machine can be removed.

L2 Cache

Currently, the cache module contains the cache implementation, datapath and controller, plus the line adapter. In order to support a 2 level cache, we will first create a cache_unit module which will just contain just the cache datapath and controller. It will have the following interface:

```
module cache_unit (
    input clk,
    input rst,

    /* Line adapter/CPU memory signals */
    input logic      mem_read,
    input logic      mem_write,
    input logic [31:0] mem_byte_enable,
    input logic [31:0] mem_address,
    input logic [255:0] mem_wdata,
    output logic      mem_resp,
    output logic [255:0] mem_rdata,

    /* Physical memory signals */
    output logic [31:0] pmem_address,
    input  logic [255:0] pmem_rdata,
    output logic [255:0] pmem_wdata,
    output logic      pmem_read,
    output logic      pmem_write,
    input  logic      pmem_resp
);
```

Then, a new top level cache_l2.sv module will instantiate two cache_unit modules as well as the line adapter. The line adapter will connect to the L1 cache and the CPU as in the current implementation, and the L2 cache will connect to the L1 cache and physical memory. The L2 byte enable input will be hard-coded to '1'.

Parameterized Cache

We will parameterize the number of sets and the number of ways in our implementation.

Sets

In order to parameterize the number of sets, the number of bits used for indexing, and therefore the number of bits in the tag, must be adjustable. To do this, the main cache module has a parameter `s_index`, which represents the number of bits used to represent the index for a given address. There are therefore 2^{s_index} sets. This is supported by parameterizing the `cache_datapath`, the `data_array`, and `array` modules. The `mem_address` must be indexed correctly in the datapath to send the different arrays the correct read and write indexes, and the array modules need to instantiate internal data of the correct size.

Ways

In order to parameterize the number of ways, we will create a module `pseudo_lru`. This module will have the following interface.

```
module pseudo_lru #(
    parameter way_bits = 1,
    parameter s_index = 3,
    parameter num_ways = 2**way_bits,
    parameter num_sets = 2**s_index
)
(
    input clk,
    input rst,
    input read,
    input load,
    input [s_index-1:0] rindex,
    input [s_index-1:0] windex,
    input [way_bits-1:0] datain,
    output [way_bits-1:0] dataout
);
```

We will keep track of the LRU state with the following structure:

```
logic [num_ways-2:0] data [num_sets-1:0]
```

Then, the logic for operating on the LRU state is as follows.

Getting LRU:

```
dataout = 0;
int lru_idx = 0;
```

```

for (int i = way_bits - 1; i >=0; i--)
    dataout[i] = ~data[rindex][lru_idx];
    lru_idx = lru_idx + lru_idx + 1 + ~data[rindex][lru_idx];

```

Setting most recently used:

```

if (load)
    int lru_idx = 0;
    for (int i = way_bits - 1; i >=0; i--)
        data[windex][lru_idx] = datain[i];
        lru_idx = lru_idx + lru_idx + 1 + datain[i];

```

In addition to the new least-recently-used logic, it is necessary to instantiate the correct number of data, dirty, valid, and tag arrays and adjust the controller logic to set load signals for these arrays appropriately. To do this, the arrays can be instantiated in a loop in a generate block.

4-Way/8-Way Set Associative Cache

Since the number of ways is already parameterized, creating a 4 or 8-way set associative cache just requires setting the parameter accordingly.

Victim Cache

The victim cache (VC) will sit in between the L1 and L2 cache. It will be a fully associative cache with 4 ways (each with 1 set). The number of ways might be incremented to 8 or 16 depending on the performance/power results. The parameterized ways/sets of advanced features will be used to accomplish this.

Initially, the data in the L1 cache will be checked. If there is an L1 miss, then the VC will be checked. Both checks happen in the same clock cycle. If there is a VC hit, then the data is moved into L1 and returned to the CPU. The evicted value from the L1 cache (if applicable) is written into the victim cache. If there is a miss in L1 and VC, then the data is read from the L2 cache. The read data is placed in the L1 cache, and the L1 evicted data (if applicable) is written into the VC cache. If an eviction from the VC occurs, it will be written to the L2 cache.

Controller Description

Following is a list of all the control signals and their default outputs.

load_tag = 1'b0

load_valid = 1'b0

load_dirty = 1'b0

load_data = 1'b0

load_lru = 1'b0

load_data_idx = 0

load_dirty_idx = 0

write_en_mux_sel = 1'b0

mem_resp = 1'b0

addr_mux_sel = 1'b0

pmem_read = 1'b0

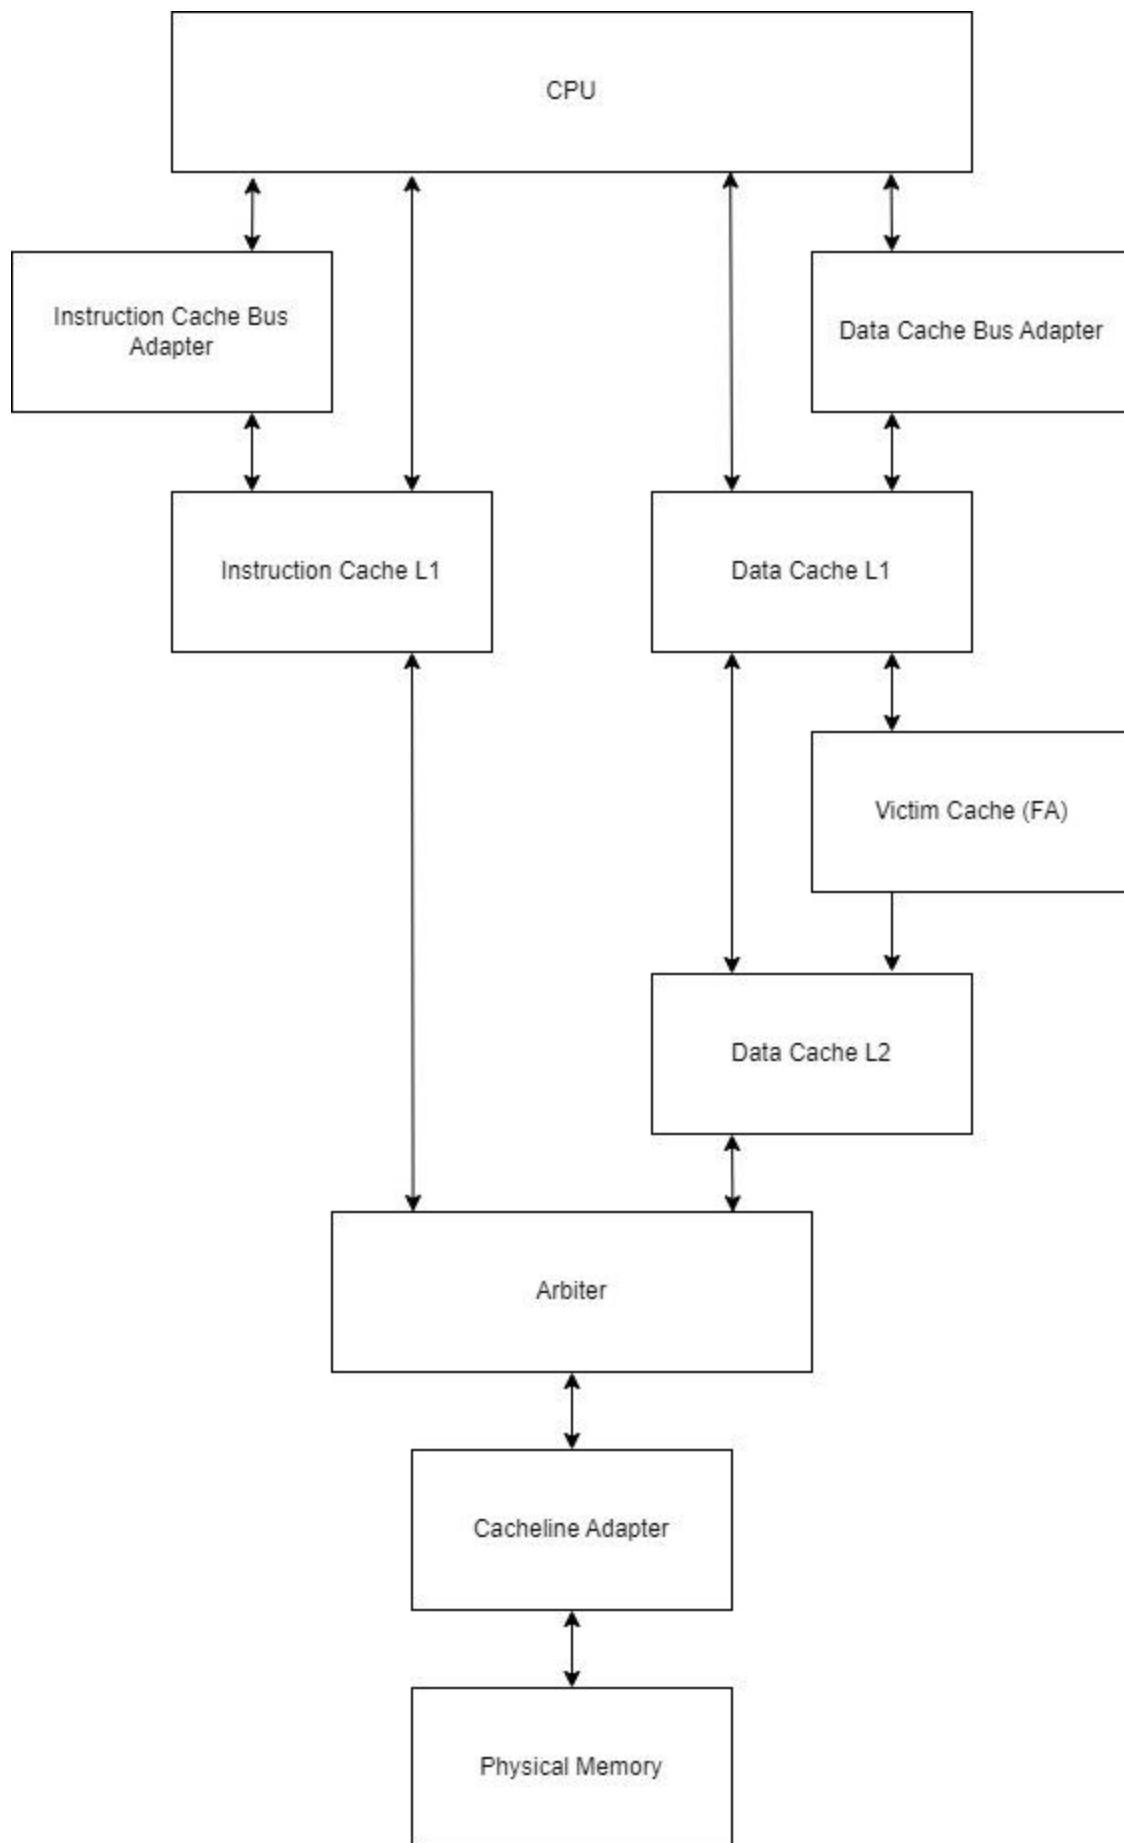
pmem_write = 1'b0

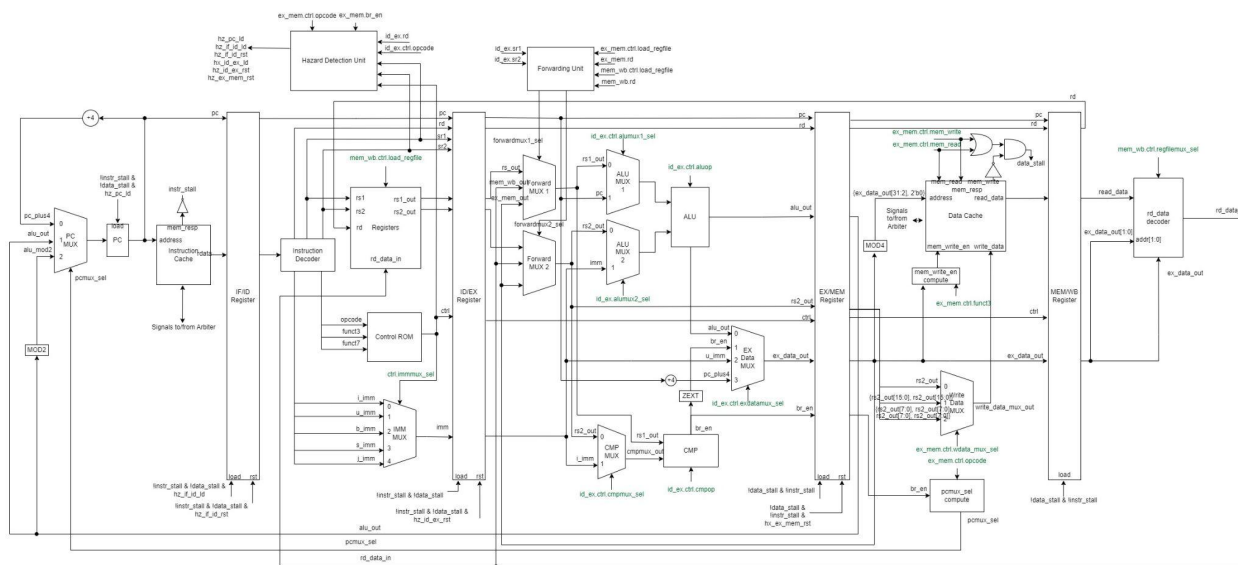

data_out_mux_sel = 1'b0

data_in_mux_sel = 1'b0

State Name	Description	Transition Conditions	Outputs
check_cache	Check for cache hit. Update LRU, and if necessary, update data, dirty, and tag arrays. This state is also an idle state.	if ((mem_read ^ mem_write) && !hit && dirty) -> write_mem; else if ((mem_read ^ mem_write) && !hit && !dirty) -> read_mem; else -> check_cache	mem_resp = 1'b1 load_lru = 1'b1 data_out_mux_sel = hit_idx if (mem_write) load_dirty = 1'b1 load_dirty_idx = hit_idx load_data = 1'b1 load_data_idx = hit_idx
write_memory	Write dirty value from cache into main memory.	if (pmem_resp) -> read_memory else -> write_memory	pmem_write = 1'b1 data_out_mux_sel = lru_idx load_dirty = 1'b1 load_dirty_idx = lru_idx

read_memory	Read from main memory and write value to cache.	if (!pmem_resp) -> read_memory else -> update_cache	pmem_read = 1'b1; addr_mux_sel = 1'b1; load_tag = 1'b1; load_valid = 1'b1; load_data = 1'b1; load_data_idx = lru_idx data_in_mux_sel = data_in_mux::rdata; write_en_mux_sel = 1'b1;
-------------	---	---	---





Branch Advanced Features:

Local Branch History Table:

The local branch history table consists of a 2-bit predictor indexed by PC values. The table contains 512 entries, with each entry containing a 2-bit predictor:

00: Strongly not taken; 01: Weakly not taken; 10: Weakly taken; 11: Strongly taken

Since we have 512 entries and all the PC addresses are 4-byte aligned, we use PC[10:2] as the table's index. Whenever there's a branch instruction, we use the PC value in the fetch stage indexing into the table to predict whether the branch is taken.

An FSM is used to update the predictor once there's a mismatch between the predicted direction and the actual direction. The actual direction and actual target address are available at MEM stage and the actual direction is the input to the FSM.

A Branch target buffer stores the target address for each branch instruction. The BTB acts like a fully-associative cache. Since the predictor table has 512 entries, the BTB also has 512 entries. In each BTB entry, we have a valid bit, a tag, and a target address. The tag is the branch PC address[31:2] because they are 4-byte aligned.

For a branch instruction, the upper 30 bits pass to the BTB and compare with each tag value to see if there's a hit. If there's a hit and the local branch history table says taken, the target address is passed to the PCmux as our next PC. If there's a hit but the local branch history table says not taken, the next PC would be PC+4. If there's a miss in BTB, the next PC would be PC+4. When the BTB encounters a miss and it's a branch instruction, we want to update the BTB with the PC address and the calculated target address. If there's an entry with a valid bit 0, fill the entry with the new data. If the buffer is fully occupied, use pseudo LRU to replace an entry with the new data.

Global 2-level Branch History Table:

The global 2-level branch history table is a shift register that records the direction of the previous branch. It's a 9-bit register, so it records 9 previous branch directions. It also has a table of the 2-bit predictor. The value in the shift register XOR with the PC [10:2] to form the index into the table, and the rest is the same as the local branch history table.

Tournament Branch Predictor:

The tournament branch predictor selects between the local branch direction output and the global branch direction output. A 2-bit FSM is used to select between the two directions. Based on the accuracy of previous local and global predictions, the FSM will update its selection bias.

00: Use_local_predictor_1; 01: Use_local_predictor_2; 10: Use_global_predictor_1; 11: Use_global_predictor_2

